

# Session 2: Introduction to R Programming

R for Stata Users

---

Luiza Andrade, Rob Marty, Rony Rodriguez-Ramirez, Luis Eduardo San Martin, Leonardo Viotti

The World Bank | [WB Github](#)

November 2021



# Table of contents

1. Introduction
2. Initial settings
3. File paths
4. Exploring a dataset
5. Creating a document outline in RStudio
6. Using packages
7. Functions inception
8. Mapping and iterations
9. Custom functions
10. Indentation
11. Appendix

# Introduction

---

# What this session is about

- In the first session, you learned how to work with R
- You are probably eager to start programming in R by now
- But before you start, we recommend learning how to write R code that will be **reproducible, efficient, intelligible and easy to navigate**
- Indeed, that's what this session is about!

# What this session is about

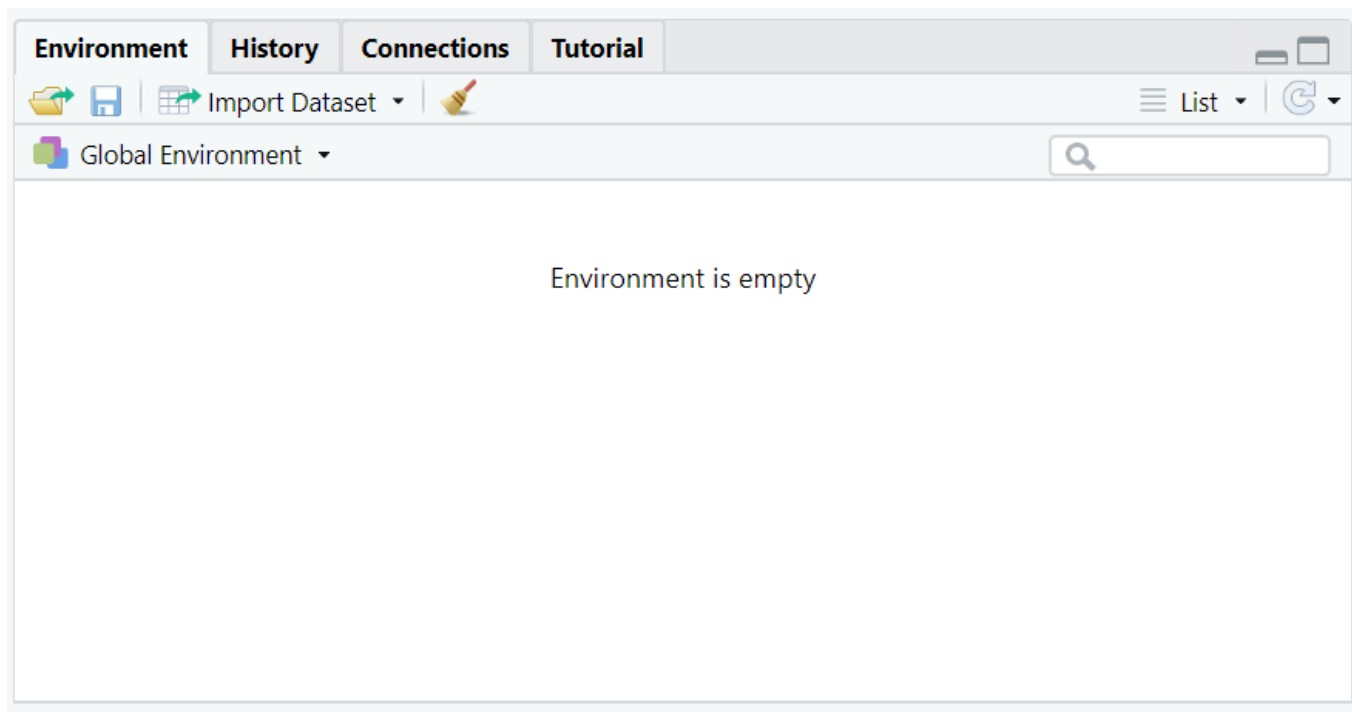
- We will cover common coding practices in R so that you can make **the most efficient use** for it
- We will also discuss some styling conventions to make your code **readable and reproducible**
- This will give you a solid foundation to write code in R and hopefully you'll be able to skip some painful steps of the "getting-your-hands-dirty" learning approach

# Initial settings

---










# Initial settings

- Let's start by opening RStudio or by closing and opening it again
- Notice two things:
  1. Your environment is *probably* empty (it's OK if it's not)



# Initial settings

- Let's start by opening RStudio or by closing and opening it again
- Notice two things:
  1. Your environment is *probably* empty (it's OK if it's not)
  2. Go to the **Console** panel and use the up and down keys to navigate through previously executed commands. They are saved by default in a file named **.Rhistory** that you might have noticed

Name	Date modified	Type
 .git	4/6/2021 2:07 PM	File folder
 .Rproj.user	4/6/2021 9:51 AM	File folder
 DataWork	4/5/2021 4:37 PM	File folder
 Presentations	4/6/2021 5:16 PM	File folder
 .gitignore	4/5/2021 4:37 PM	GITIGNORE File
 <b>.Rhistory</b>	4/6/2021 4:18 PM	RHISTORY File
 dime-r-training.Rproj	4/6/2021 4:17 PM	R Project
 LICENSE	12/15/2020 2:53 PM	File
 README.md	4/5/2021 4:37 PM	MD File

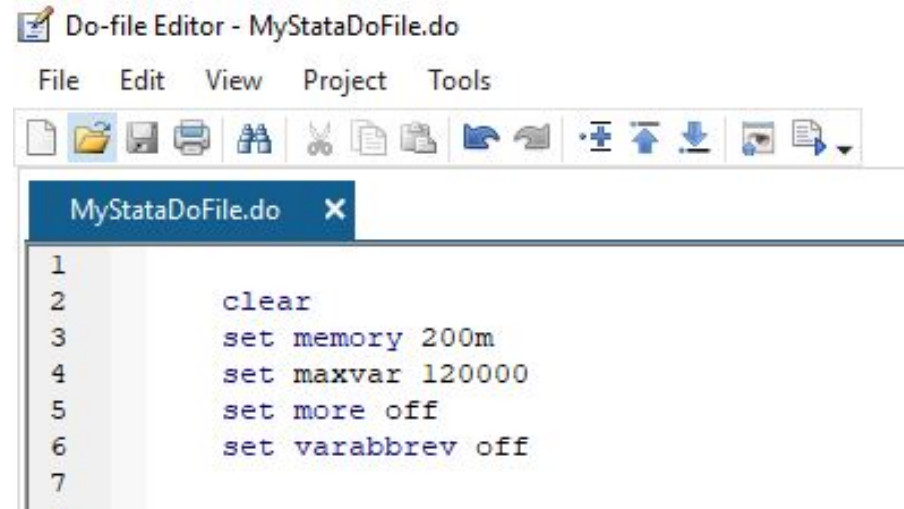


# Initial settings

- Let's start by opening RStudio or by closing and opening it again
- Notice two things:
  1. Your environment is *probably* empty (it's OK if it's not)
  2. Go to the `Console` panel and use the up and down keys to navigate through previously executed commands. They are saved by default in a file named `.Rhistory` that you might have noticed
- We'd usually want these two things -- an **empty environment** and the **history of commands** executed in previous sessions -- to be present every time we open a new RStudio session

# Initial settings

Have you ever seen these lines of code before?

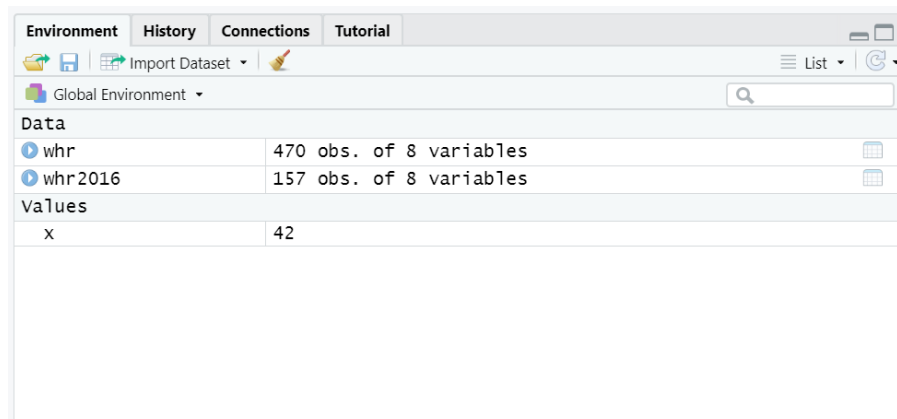


The screenshot shows the Stata Do-file Editor window titled "Do-file Editor - MyStataDoFile.do". The menu bar includes "File", "Edit", "View", "Project", and "Tools". The toolbar contains icons for opening files, saving, printing, zooming, and other editing functions. The main text area shows a do-file named "MyStataDoFile.do" with the following code:

```
1  
2      clear  
3      set memory 200m  
4      set maxvar 120000  
5      set more off  
6      set varabbrev off  
7
```

# R Initial settings

- We **don't need to set the memory or the maximum number of variables** in R
- The equivalent of `set more off` is the default
- The equivalent of `clear all` is not a default setting, but we'll change that in exercise 1
- In any case, remember that you can see all the objects in your computer's memory at any point in the `Environment` panel



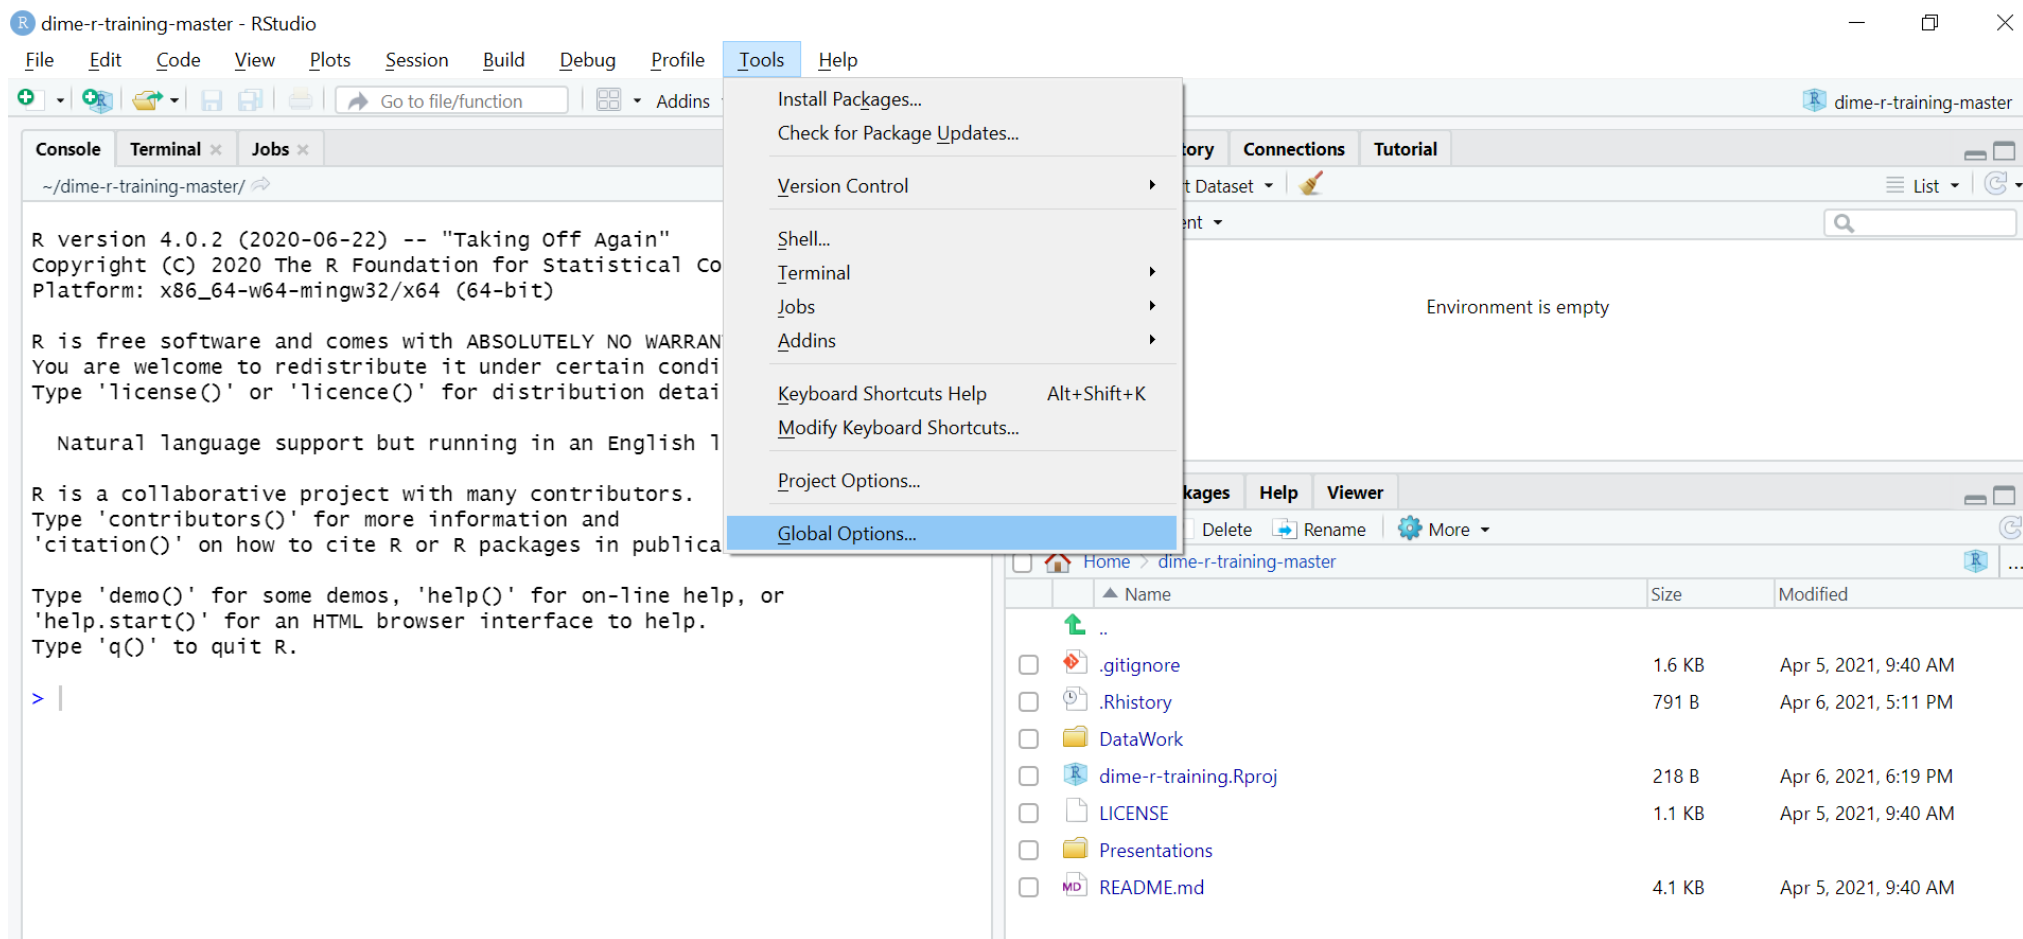
# Initial settings

## Exercise 1

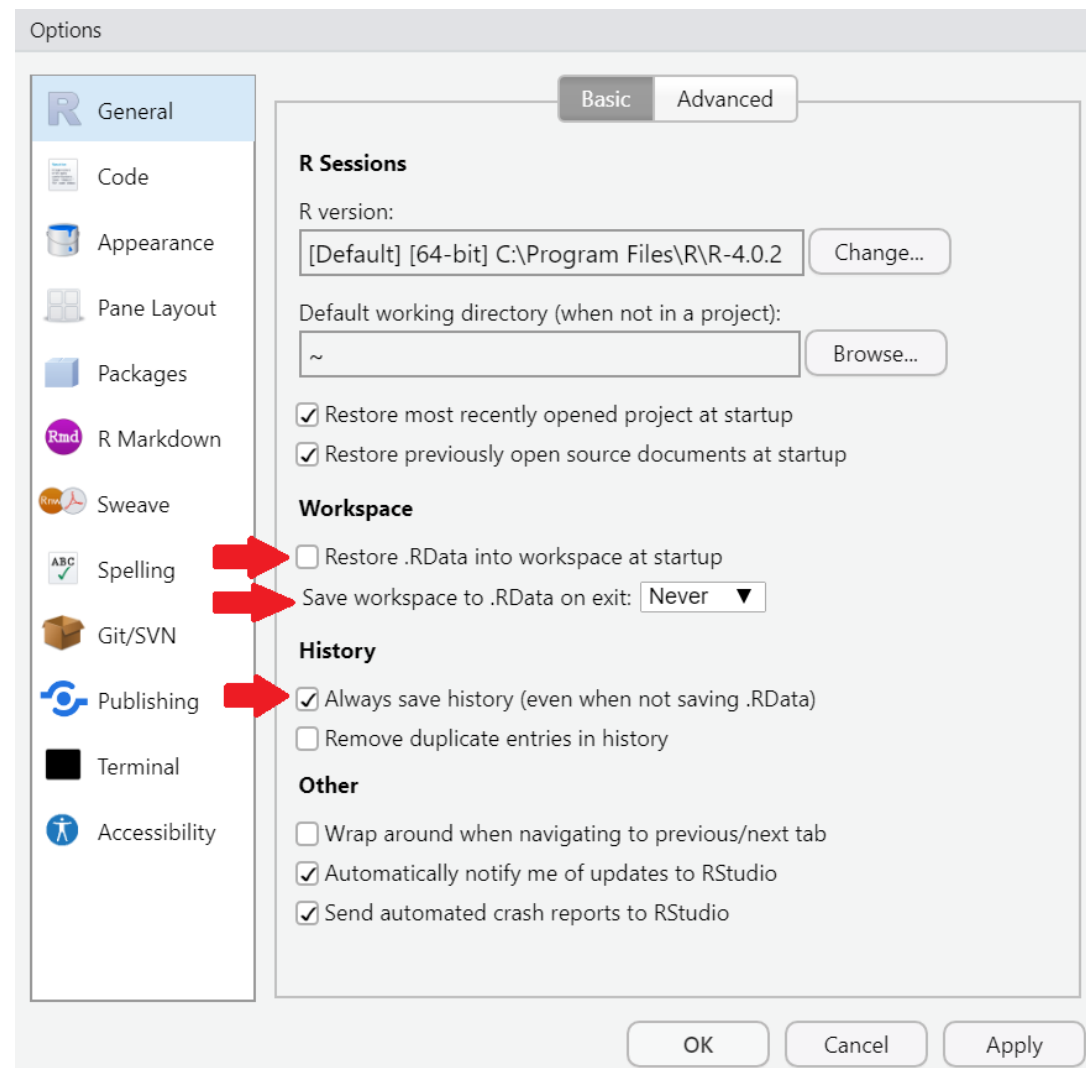
After this, you'll never have to use the equivalent of `clear all`

1. Go to `Tools` > `Global Options...`
2. In the `General` tab, make sure the following options are set:
  - Un-check *Restore .RData into workspace at startup*
  - For *Save workspace to .RData on exit*, select *Never*
  - Make sure *Always save history (even when not saving .RData)* is checked
3. Now restart RStudio

# Initial settings



# Initial settings

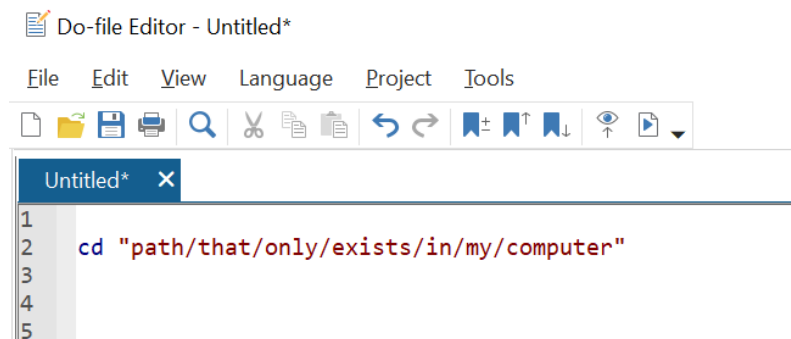


# File paths

---

# File paths

- What about working directories? We usually do something like this every time we start a new script in Stata:



- The direct equivalent to `cd` in R is this command:

```
setwd("your/path")
```

- However, we recommend not using it unless it's absolutely necessary (never, if possible)



# RStudio projects

- Instead, you should use RStudio projects and the `here` library

**Important:** We won't get into the specifics of directory organization here, but we'll assume that all the files you use for a specific project (data, scripts, and outputs) reside in the same project directory. We'll call this the **working directory**.

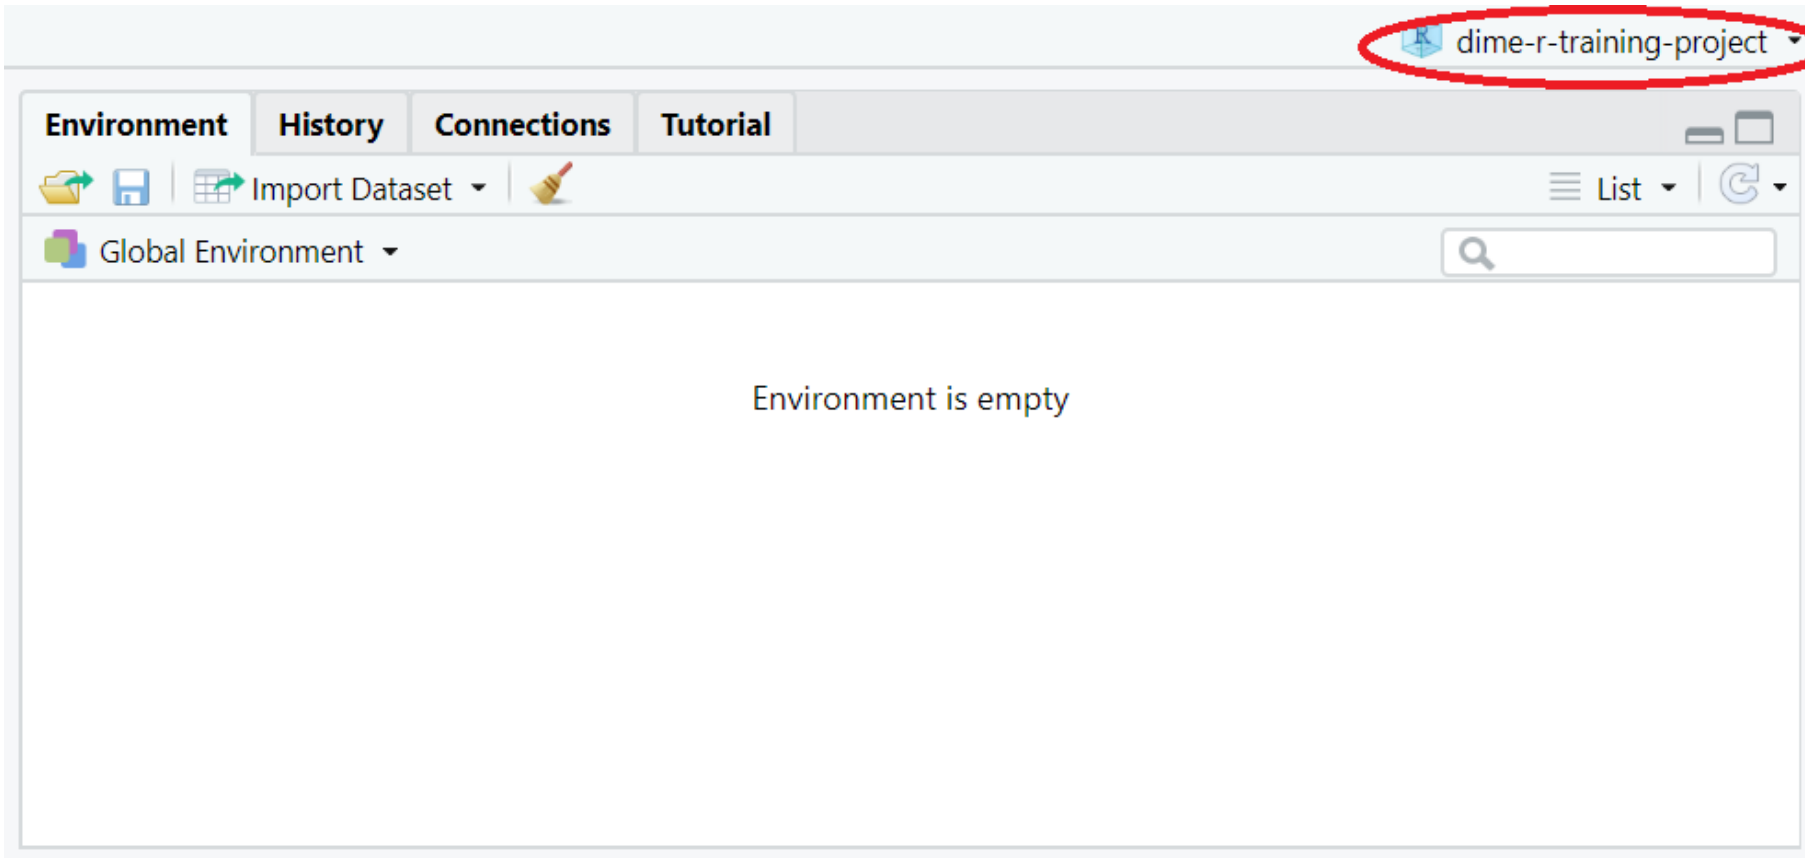
- RStudio projects let you "bind" your project files to a root directory, regardless of the path to it
- This is crucial because it allows smooth interoperability between different computers where the exact path to the project root directory differs
- Additionally, each RStudio project you work on keeps their own history of commands!

# RStudio projects

## Exercise 2

1. On RStudio, select `File` > `New Project...`
2. Select `New Directory` > `New Project`
3. Assign the name: `dime-r-training-project` to the project

# RStudio projects



# The `here` library

- `here` locates files relative to your project root
- It uses the root project directory to build paths to files easily
- Similar to RStudio projects, it allows for interoperability between different computers where the absolute path to the same file is not the same

# Usage of `here`

- Load the `here` library:

```
library(here)
```

- Now you'll be able to use `here()` to point the location of every file relative to your project root
  - For example, to load a `csv` file located in: `C:/WBG/project-root-name/data/raw/data-file.csv`, you should use:

```
path <- here("data", "raw", "data-file.csv")  
df <- read.csv(path)
```

- **Notes:**

- Your project root is the directory that contains the `.Rproj` file (or `.git` directory if you are using git).
- The result of `here()` is an absolute path that points to a file or folder location in your computer

# File paths

## Exercise 3

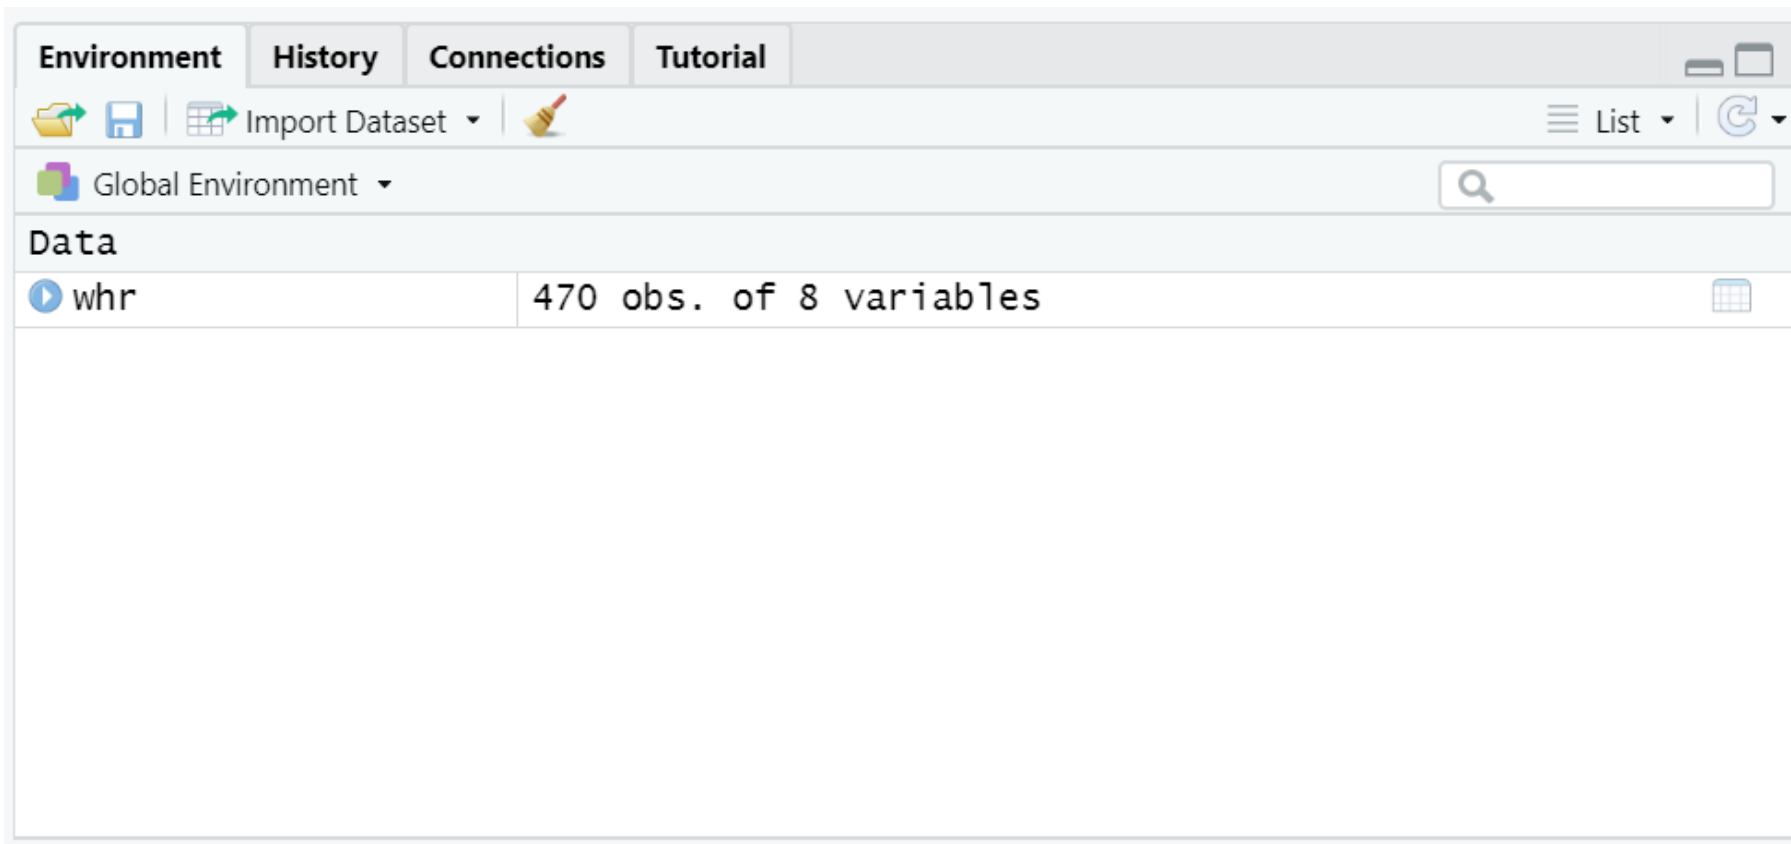
1. Go to the [OSF page of the course](#) and download the file in: `R for Stata Users - November 2021 > Data > DataWork.zip`
2. Unzip the file in your RStudio project root folder. This is the folder where the file `dime-r-training-project.Rproj` sits
3. On RStudio, go to `File > New File > R Script` and save this new empty script in `DataWork > Code` as `exercises-session2.R`
4. Now let's test if that worked. Load the `here` library and read the `csv` file `DataWork/Datasets/Final/whr_panel.csv` using `here()`
  - Use the function `read.csv()` to load the file. The argument for `read.csv()` is the result of `here()`
  - Remember to assign the dataframe you're reading to an object. You can call it `whr` as we did yesterday

# File paths

```
library(here)  
whr <- read.csv(here("DataWork", "DataSets", "Final", "whr_panel.csv"))
```

# RStudio projects and `here`

If you did the exercise correctly, you should see the `whr` data frame listed in the Environment panel





# Creating a document outline in RStudio

---

# Document outline

- RStudio allows you to **create an interactive index** for your scripts
- To add a section to your code, create a commented line with the title of your section and add at least 4 trailing dashes (`---`), pound signs (`####`) or equal signs (`====`) after it


# Document outline

## Exercise 4

1. In your script, add a header before the line where you used `library(here)` with the text: `# Part 1: Loading libraries----`
2. Add the following header before `read.csv(...)`: `Part 2: Loading data----`
  - Remember: you create a section header by adding at least 4 trailing dashes (`-`), pound (`#`) or equal (`=`) signs in a comment line
3. Note that once you create a section header, an arrow appears right next to the row number. Click on the arrows to see what happens.

# Document outline

- The outline can be accessed by clicking on the button on the top right corner of the script window. You can use it to jump from one section to another
- You can also use the keyboard shortcuts **Alt + L** (**Cmd + Option + L** on Mac) and **Alt + Shift + L** to collapse and expand sections



# Exploring a dataset

---

# Exploring a dataset

Some useful functions:

- `View()`: opens the data set
- `class()`: reports object type or type of data stored
- `dim()`: reports the size of each one of an object's dimension
- `names()`: returns the variable names of a data set
- `str()`: general information about the structure of an R object
- `summary()`: summary information about the variables in a data frame
- `head()`: shows the first few observations in the dataset
- `tail()`: shows the last few observations in the dataset

# Exploring a dataset

```
# View the data set (notice the uppercase "V")  
View(wht)
```

This is the same as clicking on the object name in the environment panel. It opens a spreadsheet-style data viewer of a dataframe.





	country	region	year	happiness_rank	happiness_score	economy_gdp_per_capita	head
1	Switzerland	Western Europe	2015	1	7.587	1.39651	
2	Iceland	Western Europe	2015	2	7.561	1.30232	
3	Denmark	Western Europe	2015	3	7.527	1.32548	
4	Norway	Western Europe	2015	4	7.522	1.45900	
5	Canada	North America	2015	5	7.427	1.32629	
6	Finland	Western Europe	2015	6	7.406	1.29025	
7	Netherlands	Western Europe	2015	7	7.378	1.32944	
8	Sweden	Western Europe	2015	8	7.364	1.33171	
9	New Zealand	Australia and New Zealand	2015	9	7.286	1.25018	

# Exploring a dataset

```
# Object type and dimensions  
class(whr)  
dim(whr)
```

Code

 Start Over

 Run Code


```
1  
2  
3
```




# Exploring a dataset

```
# Object structure  
str(whr)
```

Code

 Start Over


 Run Code


1  
2  
3

# Exploring a dataset

```
# Summarize a dataframe  
summary(whr)
```

Code

 Start Over


 Run Code


1  
2  
3

# Exploring a dataset

```
# Printing the first rows of a dataframe  
head(wht)
```

Code

 Start Over

 Run Code

```
1  
2  
3
```

# Using packages

---

# Packages

- Since there is a lot of people developing for R, it can have many different functionalities.
- To make it simpler, these functionalities are bundled into packages.
- A package is just **a unit of shareable code**.

# Packages

- It may contain new functions, but also more complex functionalities, such as a Graphic User Interface (GUI) or settings for parallel processing (similar to Stata MP).
- They can be shared through R's official repository - CRAN (13,000+ packages reviewed and tested).
- There are many other online sources such as GitHub, but it's important to be careful, as these probably haven't gone through a review process as rigorous as those in CRAN.

# Packages

- To install and use packages you can either do it with the user interface or by the command prompt.

```
# Installing a package  
install.packages("tidyverse",  
                 dependencies = TRUE)  
# the dependencies argument also installs all other packages  
# that it may depend upon to run
```

- You only have to install a package once, but you have to **load them every new session**.

# Using packages

## Exercise 5

1. Load the `tidyverse` meta package in part 1 of your script using `library(tidyverse)`
2. Run your script



# Warnings vs errors

What if this happens?

```
> library(tidyverse)
-- Attaching packages ----- tidyverse 1.3.0 --
v ggplot2 3.3.2      v purrr   0.3.4
v tibble  3.1.0      v dplyr   1.0.5
v tidyr   1.1.2      v stringr 1.4.0
v readr   1.3.1      v forcats 0.5.0
-- Conflicts ----- tidyverse_conflicts() --
x dplyr::filter() masks stats::filter()
x dplyr::lag()     masks stats::lag()

> # Part 3: Loading data----
> whr <- read.csv(here("DataWork", "DataSets", "Final", "whr_panel.csv"))
Warning messages:
1: package 'tidyverse' was built under R version 4.0.3
2: package 'tibble' was built under R version 4.0.4
3: package 'dplyr' was built under R version 4.0.4
```

# Warnings vs errors

R has two types of error messages, `warnings` and actual `errors`:

- `Errors` - break your code, usually preventing it from running.
- `Warnings` - usually mean that nothing went wrong yet, but you should be careful.

RStudio's default is to print warning messages, but not to stop the code at the lines where they occur. You can configure R to stop at warnings if you want.

# Functions inception

---

# A function inside a function

- In R, you can **write one function inside another**
- In fact, you have already done this several times in this course
- Here's an example:

# A function inside a function

```
# Print the summary of the logarithm of the happiness score

## The long way:
log_score <- log(wht$happiness_score)
summary(log_score)

# The shortcut
summary(log(wht$happiness_score))
```

# A function inside a function

- This is a simple example of **metaprogramming** (that's the real name of this technique) and may seem trivial, but it's not
- For starters, you can't do it in Stata!

# A function inside a function

```
(R)
 _ _ _ _ _
|_|_|_|_|_|
|_|_|_|_|_|
|_|_|_|_|_|
|_|_|_|_|_|
|_|_|_|_|_|
Statistics/Data Analysis

MP - Parallel Edition

15.1 Copyright 1985-2017 StataCorp LLC
      StataCorp
      4905 Lakeway Drive
      College Station, Texas 77845 USA
      800-STATA-PC      http://www.stata.com
      979-696-4600      stata@stata.com
      979-696-4601 (fax)

681-user 4-core Stata network perpetual license:
      Serial number: 501506002486
      Licensed to: WBG User
                  World Bank Group

Notes:
  1. Unicode is supported; see help unicode\_advice.
  2. More than 2 billion observations are allowed; see help obs\_advice.
  3. Maximum number of variables is set to 120000; see help set\_maxvar.
  4. New update available; type -update all-

running C:\Program Files (x86)\Stata15\sysprofile.do ...

. sysuse auto
  (1978 Automobile Data)

. summarize log(make)
variable log not found
r(111);

.
```

# A function inside a function

- Metaprogramming is a **very powerful technique**, as you will soon see
- It's **also a common source of error**, as you can only use one function inside the other if the output of the inner function is the same as the input of the outer function
- It can also get quite tricky to follow what a line of code with multiple functions inceptions is doing



# Piping

- Ever heard of piping? It's this: `%>%`
- Piping is a way of doing metaprogramming
- The actual meaning of the pipes is: *Pipes take the **output** of the function at the left and pass it as the **first argument** of the function at the right*
- The advantages of using piping is that it allows to have a cleaner division of successively applied functions in R code, **drastically improving code readability**

# Piping

```
# 1: Doing it the long way -----  
log_score <- log(whr$happy_score)  
mean(log_score)  
  
# 2: Shortcut to get to the same place -----  
mean(log(whr$happy_score))  
  
# 3: Now with pipes -----  
whr$happy_score %>%  
  log() %>%  
  mean()
```

Just remember:

- `x %>% f()` is the same as `f(x)`
- `x %>% f() %>% g()` is the same as `g(f(x))`

# Mapping and iterations

---

# Iterations in R

- In Stata, we use `for` loops very frequently
- In R, the syntax of `for` loops is this:

```
for (number in 1:3) {  
  print(number)  
}
```

```
## [1] 1  
## [1] 2  
## [1] 3
```

- R, however, has a set of functions that allows users to loop through an object **in a more efficient way**, without using explicit loops
- In this training we'll introduce `map()`. It is a function part of the tidyverse meta package
- Also, in case you have not noticed yet: **R is vectorized!** this means that many operations are applied element-wise by default so you don't have to code loops to apply them to each element of a vector or dataframe

# Map

- To use `map()`, you need to either load the package `purrr` either by itself or alongside other `tidyverse` packages
- The basic syntax of `map()` is:


`map(X, function, ...)`: applies `function` to each of the elements of `X`. If `X` is a data frame then `function` is applied column-wise while if it's a vector or a list it is applied item-wise. The output of `map()` is always a vector with the results.


- **X**: a data frame, matrix or vector the function will be applied to
- **function**: the function you want to apply to each of the elements of `X`

# Map

```
# Round the values of the following vector  
x <- c(1.2, 2.5, 9.1, 5.8)  
x %>% map(round) # Rounding the vector elements  
round(x)         # since R is vectorized, this also works
```

Code

 Start Over

 Run Code

```
1  
2  
3
```

# Map vs looping

- When looping, you repeat the same operation over a set of items
- `map()`, instead, takes all your elements at once and applies an operation to them simultaneously
- The difference is like this:
  - Imagine you ask a yes/no question to a group of people
  - You can collect the answers by asking each one of them individually -- **this is looping**
  - Otherwise, you can ask them to raise their hands and collect all answers at once -- **this is** `map()`
- The output of a loop is the regular output of the operation you're repeating, times the number of iterations you did
- The output of `map()` will be always a list

# Map

## Exercise 6

1. Select the columns `freedom` and `happiness_score` of `whr`
  - Use tidyverse's `select()` for this, as in: `whr %>% select(freedom, happiness_score)`
2. Use `map()` to get the maximum value of these columns

Hints:


- Remember the syntax of `map()`: `map(X, function)`
- Try using pipe concatenation here
- The function to get the maximum over a set of numbers is `max`.




# Map

```
whr %>%  
  select(freedom, happiness_score) %>%  
  map(max)
```

Code

 Start Over

 Run Code

1  
2  
3

# List indexing revisited: `[]` vs `[][]`


- The output of `map()` will be always a list
- Remember we said yesterday that you could use brackets to index list items? we can also use double brackets
- While both will return the i-th element, a key difference is that they return different data types:
  - `[]`: returns a list of one item containing the element selected
  - `[][]`: returns the element selected itself
- This difference is crucial when working with `map()`

# List indexing revisited: `[]` vs `[[ ]]`

```
my_list <- list(1, "hello")  
my_list[2]    # this returns the string "hello" enclosed in another list  
my_list[[2]]  # this returns the actual string "hello"
```

Code

 Start Over

 Run Code

```
1  
2  
3
```

# Map

- `map()` looks nice, doesn't it?
- But what about cases when it's impossible to implement the operations I want to apply in only one function? Do I have to use `for` loops then?
- Not at all! Let's get to the next section for those cases.

# Custom functions

---


# Writing your own functions


- As we have said several times, **R is super flexible**
- One example of that is that it's **super easy and quick to create custom functions**
- Here's how:

# Custom functions

```
square <- function(x) {  
  y <- x ^ 2  
  return(y)  
}
```

Code

 Start Over

 Run Code

1  
2  
3

# Custom functions

## Exercise 7

Create a function named `zscore` that standardizes the values of a vector.

### Hints:

- The command to obtain the mean of a vector is `mean(x)`
- The command to get the SD of a vector is `sd(x)`
- R is vectorized: you can operate vectors and numbers directly and the result will be a vector
- Don't forget to include the argument `na.rm = TRUE` in `mean()` and `sd()`
- Recall the syntax of custom functions in R:

```
function_name <- function(input) {  
  
  output <- operation(input)  
  
  return(output)  
  
}
```



# Custom functions

```
zscore <- function(x) {  
  mean <- mean(x, na.rm = TRUE)  
  sd    <- sd(x, na.rm = TRUE)  
  z     <- (x - mean)/sd  
  return(z)  
}
```

# Custom functions

## Exercise 8

1. Use tidyverse's `select()` to select the columns `health_life_expectancy` and `freedom` in `whr`
2. Use `map()` combined with the `zscore` function to get the z-score of these two columns and assign the resulting list to an object named `z_scores`
3. Use list indexing on `z_scores` to generate two new columns in `whr` with the standardized values of `health_life_expectancy` and `freedom`

### Hints:

- Use the pipes (`%>%`) successively
- Remember that we don't use parenthesis next to the function name we're using `map()` with
- Remember that we use double brackets instead of single brackets to index the actual elements of a list

# Custom functions

```
z_scores <- whr %>%  
  select(health_life_expectancy, freedom) %>%  
  map(zscore)  
whr$hle_st <- z_scores[[1]]  
whr$freedom_st <- z_scores[[2]]
```

# Indentation

---

# Indentation

*# Here's some code*


```
annualHappy_reg <- aggregate(happy_score ~ year + region, data = whr, FUN = mean)
ggplot(annualHappy_reg, aes(y = happy_score, x = as.factor(year), color = region, group = region)) +
  geom_line() + geom_point()
```

*# Here's the same code*

```
annualHappy_reg <-
  aggregate(happiness_score ~ year + region,
            data = whr,
            FUN = mean)

ggplot(annualHappy_reg,
       aes(y = happiness_score,
           x = as.factor(year),
           color = region,
           group = region)) +
  geom_line() +
  geom_point()
```

# Indentation



# Why indent?

- R understands what unindented code says, but it can be **quite difficult for a human being to read it**
- On the other hand, **white space does not have a special meaning for R**, so it will understand code that is more readable for a human being

# Why indent?

- Indentation in R looks different than in Stata:
  - To indent a whole line, you can select that line and press **Tab**
  - To unindent a whole line, you can select that line and press **Shift + Tab**
  - However, this will not always work for different parts of a code in the same line
- In R, we typically don't introduce white space manually
- It's rather introduced by RStudio for us



# Indentation

## Exercise 9

To see an example of how indenting works in RStudio, let's use an example with `map()`:

```
# An elegant "loop" in R  
map(c(1.2, 2.5, 9.1, 5.8), round)
```

1. Add a line between the two arguments of the function (the vector of numbers and `round`)
2. Now add a line between the numbers in the vector.

# Indentation

Note that RStudio formats the different arguments of the function differently:

```
# A much more elegant loop in R  
map(c(1.2,  
      2.5,  
      9.1,  
      5.8),  
     round)
```

Thank you!

# Appendix

# Appendix - `.Rhistory` and `.RData`

- `.Rhistory` automatically stores the commands entered in the console
- `.RData` stores the objects in your environment only if you save your workspace, and loads them again in the next RStudio session
- Both files are relative to the working directory where your RStudio session started

# Appendix - Using packages

Once a package is loaded, you can use its features and functions. Here's a list of some useful packages:

- `Rcmdr` - easy to use GUI
- `swirl` - an interactive learning environment for R and statistics.
- `ggplot2` - beautiful and versatile graphics (the syntax is a pain, though)
- `stargazer` - awesome latex regression and summary statistics tables
- `foreign` - reads `.dta` and other formats from inferior statistical software
- `zoo` - time series and panel data manipulation useful functions
- `data.table` - some functions to deal with huge data sets
- `sp` and `rgeos` - spatial analysis
- `multiwayvcov` and `sandwich` - clustered and robust standard errors
- `RODBC`, `RMySQL`, `RPostgresSQL`, `RSQLite` - For relational databases and using SQL in R.

# Appendix - Resources

- A discussion of folder structure and data management can be found here:  
[https://dimewiki.worldbank.org/wiki/DataWork\\_Folder](https://dimewiki.worldbank.org/wiki/DataWork_Folder)
- For a broader discussion of data management, go to [https://dimewiki.worldbank.org/wiki/Data\\_Management](https://dimewiki.worldbank.org/wiki/Data_Management)

# Appendix - Git

Git is a version-control system for tracking changes in code and other text files. It is a great resource to include in your work flow.

We didn't cover it here because of time constraints, but below are some useful links, and DIME Analytics provides trainings on Git and GitHub, so keep an eye out for them.

- **DIME Analytics git page:** <https://worldbank.github.io/dimeanalytics/git/>
- **A Quick Introduction to Version Control with Git and GitHub:** <https://journals.plos.org/ploscompbiol/article?id=10.1371/journal.pcbi.1004668>



# Appendix - More on R projects

If you want to learn more about them, we recommend starting here: <https://r4ds.had.co.nz/workflow-projects.html>

# Appendix - Commenting

- To comment a line, write `#` as its first character

```
# This is a comment  
print("But this part is not")
```

- You can also add `#` halfway through a line to comment whatever comes after it

```
print("This part is not a comment") # And this is a comment
```

- In Stata, you can use `/*` and `*/` to comment in the middle of a line's code. That is not possible in R: everything that comes after `#` will always be a comment
- To comment a selection of lines, press `Ctrl + Shift + C`

# Appendix - Assignment 1

## Exercise

1. In your script panel, select all the lines of your script
2. Use the keyboard shortcut to comment these lines.
  - Shortcut: `Ctrl` + `Shift` + `C`
3. Use the keyboard shortcut to comment these lines again. What happened?

# Appendix - Looping

- In Stata, we use `for` loops very frequently
- The equivalent to that in R would be this

```
# A for loop in R  
for (number in 1:5) {  
  print(number)  
}
```

# Appendix - Column extraction operators

- Remember the use of `$` to extract columns from a dataframe?
- Other than `$`, we can also use double brackets to extract the column of a dataframe:

```
# With $:
```

```
whr$year
```

```
# With [[]]:
```

```
whr[["year"]] # Notice the use of double quotes
```

# Appendix - Column extraction operators: `[[[]]` vs `$`

What's the key difference between them?

Well, `[[[]]` lets us use other objects to refer to column names, while `$` doesn't

```
col_name <- "year"
head(whr$col_name) # this returns a NULL object because no column has the name "col_name" in whr
```

```
## Warning: Unknown or uninitialised column: `col_name`.
```

```
## NULL
```

```
col_name <- "year"
head(whr[[col_name]])
```

```
## [1] 2015 2015 2015 2015 2015 2015
```

# Appendix - Column extraction operators: `[[ ]]` vs `$`

This difference is key because we can use `[[ ]]` to loop through column names, while this is not directly possible with `$`.

```
# Printing the first observation of every column of whr
for (col in colnames(whr)) {
  whr[[col]] %>%
    head(1) %>%
    print()
}
```

```
## [1] "Switzerland"
## [1] "Western Europe"
## [1] 2015
## [1] 1
## [1] 7.587
## [1] 1.39651
## [1] 0.94143
## [1] 0.66557
```

# Appendix - Apply

- Apart from tidyverse's `map()`, base R also has a set of functions that allows users to apply a function to a number of objects without using explicit loops
- They're called `apply` and there are many of them, with different use cases
- If you look for the `apply` help file, you can see all of them
- We'll show only two of them, `sapply` and `apply`



# Appendix - Apply

- The syntax of `sapply()` is:

```
sapply(X, FUN, ...)
```

- Its main arguments are:
  - **X**: a data frame, matrix or vector the function will be applied to
  - **FUN**: the function you want to apply
- `sapply()` applies the function (`FUN`) to all the elements of `X`. If `X` is a data frame then the function is applied column-wise, while if it's a vector or a list it is applied item-wise
- The output of `sapply()` is usually a vector with the results, but it can be a matrix if the results have more than one dimension

# Appendix - Apply

```
# A for loop in R  
for (number in c(1.2, 2.5)) {  
  print(round(number))  
}
```

```
# A much more elegant option  
sapply(c(1.2, 2.5), round)
```

# Appendix - Apply

```
# Printing the first observation of every column of whr  
for (col in names(whr)) {  
  print(head(whr[[col]], 1))  
} # Option 1  
  
sapply(whr, head, 1) # A more elegant and efficient option
```

# Appendix - Apply

- A more general version of `sapply()` is the `apply()` function. This is its syntax:

```
apply(X, MARGIN, FUN, ...)
```

- Arguments:
  - **X**: a data frame (or matrix) the function will be applied to
  - **MARGIN**: 1 to apply the function to all rows or 2 to apply the function to all columns
  - **FUN**: the function you want to apply
- `apply()` applies a function (`FUN`) to all columns or rows of matrix (`X`). A value of 1 in `MARGIN` indicates that the function should be applied row-wise, while 2 indicates columns

# Appendix - Apply

```
matrix <- matrix(c(1, 24, 9, 6, 9, 4, 2, 74, 2), nrow = 3) # Defining a matrix  
apply(matrix, 1, mean) # row means  
apply(matrix, 2, mean) # column means
```

# Appendix - Assignment 2

## Exercise: Get the row max

1. Select the columns `freedom` and `happiness_score` of `whr`
  2. Use `apply()` to get the row max between these two columns, for every row
- Hints:
    - Remember the syntax of `apply()`: `apply(X, MARGIN, FUN)`
    - A value of 1 for `MARGIN` indicates that the function must be applied row-wise
    - The function to get the maximum over a set of numbers is `max`

# Appendix - Apply

Solution:

```
whr %>%  
  select(freedom, happiness_score) %>%  
  apply(1, max)
```