#### R programming practices

R for Stata Users

DIME Analytics The World Bank | WB Github April 2021



#### Table of contents



- 1. Introduction
- 2. Initial settings
- 3. File paths
- 4. Exploring a dataset
- 5. Commenting
- 6. Creating a document outline in RStudio
- 7. Using packages
- 8. Functions inception
- 9. Looping
- 10. Custom functions
- 11. Indentation
- 12. Appendix

### Introduction

#### Introduction



#### What is this session about?

- In the previous sessions, you learned how to use R in data analysis
- You are probably eager to get your hands into some data using R by now, and you would figure out **what should be in your code** for it to work
- But you would probably not know right away **how to write that**, so that in the end you might have code that is only intelligible for yourself -- and not for a very long time

#### Introduction



#### What is this session about?

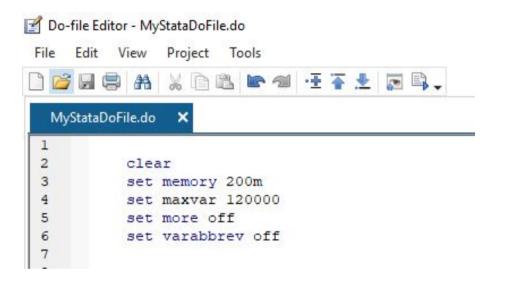
- In this session, we will cover common coding practices in R so that you can make **the most efficient use** for it
- We will also discuss some styling conventions to make your code **readable and reproducible**
- This will give you a solid foundation to code in R, and hopefully you'll be able to skip some painful steps of the "getting-your-hands-dirty" learning approach



- Let's start by opening RStudio, or by reopening it if you had it opened already
- What do you see in your environment?
- If you saved the objects from the last session in .RData, the objects that were in RStudio's memory last time you closed it will still be there whenever you open it again
- You can also go to the <code>Console</code> panel and use the up and down keys to navigate through the commands you executed in previous sessions. They are saved by default in <code>.Rhistory</code>



Raise your hand if you have ever seen these lines of code before:

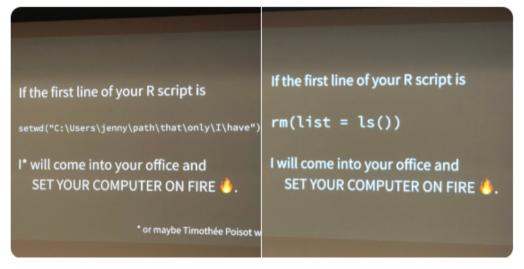




- ullet We don't need to set the memory or the maximum number of variables in R
- The equivalent of set more off is the default
- You can see all the objects in your memory at any point in the **Environment** pane







4:50 PM - 10 Dec 2017



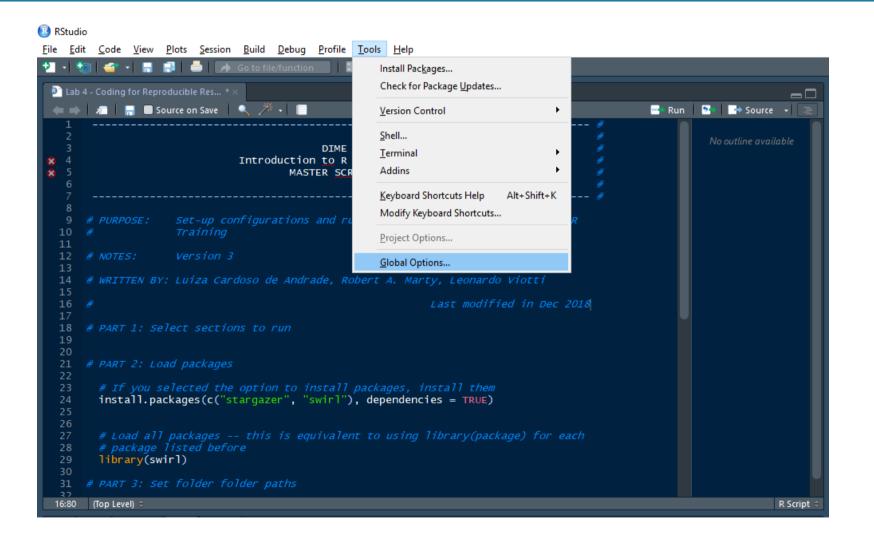


You don't need to use the equivalents to Stata's cd() and clear all in R if set global options correctly in RStudio.

# Exercise 1: To make sure that no one will burn your computer, here's how you change these settings:

- 1. Go to Tools > Global options...
- 2. Make sure the following options are checked:
  - Re-use idle sessions for project links (Note: this is only relevant for Windows)
  - Restore most recently opened project at startup
  - Restore previously open source documents at startup
  - Restore .RData into workspace at startup
  - Always save history (even when not saving .RData)
- 3. Now restart RStudio

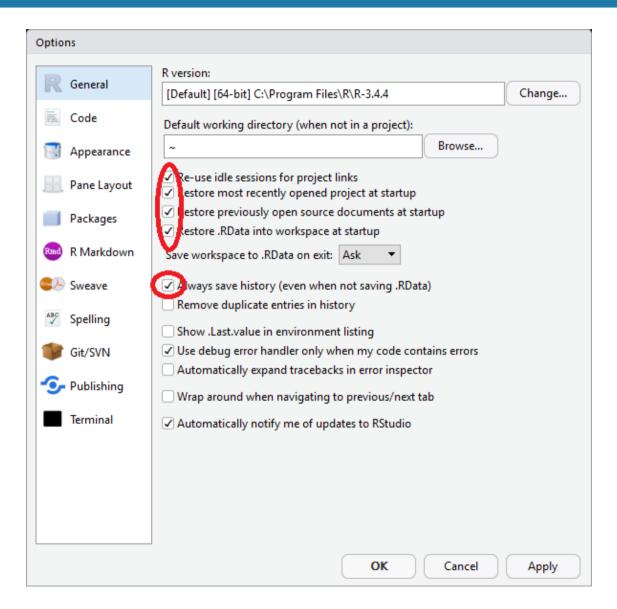






Options		
R General	R version:	
General	[Default] [64-bit] C:\Program Files\R\R-3.4.4	Change
Code	Default working directory (when not in a project):	
Appearance	~ Browse	
Pane Layout	<ul> <li>Re-use idle sessions for project links</li> <li>Restore most recently opened project at startup</li> </ul>	
Packages	<ul> <li>Restore previously open source documents at startup</li> <li>Restore .RData into workspace at startup</li> </ul>	
R Markdown	Save workspace to .RData on exit: Ask ▼	
Sweave	Always save history (even when not saving .RData)	
Spelling Spelling	Remove duplicate entries in history  Show .Last.value in environment listing	
Git/SVN	Use debug error handler only when my code contains errors	
• Publishing	Automatically expand tracebacks in error inspector      Wrap around when navigating to previous/next tab	
Terminal	✓ Automatically notify me of updates to RStudio	
	OK Cancel	Apply







- For the purpose of this session, we will assume that you are dealing with a specific folder structure.
- Folder organization, though an important part of data work, is outside the scope of this course.
- You can find resources about it in the appendix, and we have shared with you a folder that is organized as we want it to be.
- For today's session:
  - Go to the OSF repository of this course (https://osf.io/b7hm6/)
  - Download the file DataWork.zip
  - o Download and unzip the materials in a location you can remember
  - Then go to the DataWork/Code folder and open the file Lab 2 Intro II.R
- We will use this script as a basis for the exercises, and you should modify it during this session to complete them.



- In the last session, we used the menu bar to load a data set into R
- Today, we will do that using code and referring a file's path
- File paths in R, as in Stata, are basically just strings
- Note, however, that in R we can only use forward slashes (/) to separate folder names

#### Exercise 3: File path to your folder

Let's start by adding the folder path to the training's folder in your computer to the **projectFolder** object at the beginning of **PART 2** 

```
# Add your folder path here
# Remember to replace backslashes ("\") with regular ones ("/")
projectFolder <- "YOUR/FOLDER/PATH/HERE"</pre>
```



- You can set file paths in your script using the file.path() function
- This function concatenates strings using / as a separator to create file paths
- This is similar to using globals to define folder paths in Stata



Let's test if that worked:

```
# Project folder
projectFolder <-
    "/home/luis_eduardo/cs_projects/dime-r-training"

# Data work folder
dataWorkFolder <- file.path(projectFolder, "DataWork")

# Print data work folder
dataWorkFolder</pre>
```

```
## [1] "/home/luis_eduardo/cs_projects/dime-r-training/DataWork"
```



#### Loading a dataset from CSV

```
Use the read.csv() command:
```

```
read.csv(file, header = FALSE, stringsAsFactors = FALSE)
```

- file: is the path (string) to the file you want to open, including its path, name and file extension ( .csv ).
  - Example: /home/user/Data/whr\_panel.csv
- header: if TRUE, will read the first row as variable names (default is header = FALSE).
- **stringsAsFactors**: logical. See next slide for more.



#### Loading a dataset from CSV

- Since R 4.0.0 and beyond, stringsAsFactors = FALSE is the default. In every previous version, the default is TRUE.
- This means that if your R version is 3.X.X, R will turn any string variables into factors by default when reading a csv file.
- This format **saves memory**, but can be tricky if you actually want to use the variables as strings.
- You can specify the option stringsAsFactors = FALSE to make sure you prevent R from turning strings into factors.



#### Exercise 4: Test file paths

- 1. Save your script.
- 2. Start a new R session: go to Session > New Session. This session should be completely blank.
- 3. Open the code you just saved.

```
• Go to File > Open File... and select your script
```

4. Add a line opening the data set in your code

1. Run the whole script. If it worked, your environment should now include the whr dataset and the path locals.



#### Some useful functions:

- View(): open the data set
- class(): reports object type or type of data stored
- dim(): reports the size of each one of an object's dimension
- names(): returns the variable names of a data set
- str(): general information about the structure of an R object
- **summary():** summary information about the variables in a data frame
- head(): shows the first few observations in the dataset
- tail(): shows the last few observations in the dataset



#### Exercise 5: Explore a dataset

Use some of these functions to explore the whr data set.

- View()
- class()
- dim()
- names()
- str()
- summary()
- head()
- tail()



# View the data set (same as clicking on it in the Environment pane)
View(whr)



```
class(whr)

## [1] "data.frame"

dim(whr)

## [1] 470 8
```



str(whr)



#### summary(whr)

```
##
                          region
                                                         happiness rank
     country
                                               vear
    Length: 470
                       Length: 470
                                                 :2015
                                                         Min. : 1.00
##
                                          Min.
   Class :character
                       Class :character
                                          1st Qu.:2015
                                                        1st Qu.: 40.00
                                                        Median : 79.00
##
   Mode :character
                       Mode :character
                                          Median :2016
##
                                                 :2016
                                                               : 78.83
                                          Mean
                                                         Mean
##
                                          3rd Qu.:2017
                                                        3rd Qu.:118.00
##
                                                 :2017
                                                                :158.00
                                          Max.
                                                         Max.
   happiness_score economy_gdp_per_capita health_life_expectancy
                                                                     freedom
##
##
   Min.
          :2.693
                    Min.
                           :0.0000
                                           Min.
                                                  :0.0000
                                                                  Min.
                                                                         :0.0000
                                                                  1st Qu.:0.2976
   1st Qu.:4.509
                   1st Qu.:0.6053
                                           1st Qu.:0.4023
##
   Median :5.282
                                           Median : 0.6301
##
                    Median :0.9954
                                                                  Median :0.4183
##
   Mean
          :5.371
                    Mean
                           :0.9278
                                           Mean
                                                  :0.5800
                                                                  Mean
                                                                         :0.4028
   3rd Qu.:6.234
                    3rd Qu.:1.2524
                                           3rd Qu.:0.7683
                                                                  3rd Qu.:0.5169
##
##
   Max.
           :7.587
                           :1.8708
                                           Max.
                                                  :1.0252
                                                                  Max.
                                                                          :0.6697
                    Max.
```



head(whr)

```
##
                         region year happiness_rank happiness_score
         country
## 1 Switzerland Western Europe 2015
                                                               7.587
## 2
         Iceland Western Europe 2015
                                                               7.561
## 3
         Denmark Western Europe 2015
                                                               7.527
## 4
         Norway Western Europe 2015
                                                               7.522
## 5
         Canada North America 2015
                                                               7.427
## 6
         Finland Western Europe 2015
                                                               7.406
     economy_gdp_per_capita health_life_expectancy freedom
##
## 1
                    1.39651
                                            0.94143 0.66557
## 2
                    1.30232
                                            0.94784 0.62877
## 3
                    1.32548
                                            0.87464 0.64938
## 4
                                            0.88521 0.66973
                    1.45900
## 5
                    1.32629
                                            0.90563 0.63297
## 6
                    1.29025
                                            0.88911 0.64169
```



Didn't get all of those?

Don't worry, you'll see them again soon.

# Commenting

#### Commenting



- To comment a line, write # as its first character
- You can also add # half way through a line to comment whatever comes after it
- In Stata, you can use /\* and \*/ to comment part of a line's code. That is not possible in R: whatever comes after # will be a comment
- To comment a selection of lines, press Ctrl + Shift + C

### Commenting



#### Exercise 6: Commenting

- 1. Go the Lab 2 Intro II.R script. Select the lines under PART 2: Set folder paths.
- 2. Use the keyboard shortcut to comment these lines.
  - Shortcut: Ctrl + Shift + C
- 3. Use the keyboard shortcut to comment these lines again. What happened?

# Creating a document outline in RStudio

#### Creating a document outline in RStudio



- RStudio also allows you to **create an interactive index** for your scripts
- To add a section to your code, create a commented line with the title of your section and add at least 4 trailing dashes ( --- ), pound signs ( #### ) or equal signs ( ==== ) after it

### Creating a document outline in RStudio



#### Exercise 7: Headers

- 1. Open the script index and make PART 1 a section header. Do the same for parts 2 and 3.
  - Remember: you create a section header by adding at least 4 trailing dashes ( ), pound ( # ) or equal ( = ) signs in a comment line
- 2. Note that once you create a section header, an arrow appears right next to it. Click on the arrows of parts 2 and 3 to see what happens.

### Creating a document outline in RStudio



- The outline can be accessed by clicking on the button on the top right corner of the script window. You can use it to jump from one section to another
- You can also use the keyboard shortcuts Alt + L (Cmd + Option + L on Mac) and Alt + Shift + L to collapse and expand sections



#### Packages

- Since there is a lot of people developing for R, it can have many different functionalities.
- To make it simpler, these functionalities are bundled into packages.
- A package is just a unit of shareable code.



#### **Packages**

- It may contain new functions, but also more complex functionalities, such as a Graphic User Interface (GUI) or settings for parallel processing (similar to Stata MP).
- They can be shared through R's official repository CRAN (13,000+ packages reviewed and tested).
- There are many other online sources such as GitHub, but it's important to be careful, as these probably haven't gone through a review process as rigorous as those in CRAN.



#### **Packages**

• To install and use packages you can either do it with the user interface or by the command prompt.

• You only have to install a package once, but you have to load it every new session.



#### Exercise 8

1. Now load the package we just installed. Use the library() function to do it.



library(tidyverse)

Notice that we used quotes around the name to install the package, but we don't need them anymore to load it.



#### Warnings vs errors

What if this happens?



#### Warnings vs errors

R has two types of error messages, warnings and actual errors:

- Errors break your code, usually preventing it from running.
- Warnings usually mean that nothing went wrong yet, but you should be careful.

RStudio's default is to print warning messages, but not stop the code at the lines where they occur. You can configure R to stop at warnings if you want.



- In R, you can write one function inside another
- In fact, you have already done this a few times in this course
- Here's an example:



```
# Doing it the long way -----
  # Create a vector with the log of the happiness score
  log score <- log(whr$happiness score)</pre>
  # Get descriptive statistics for the log vector
  summary(log score)
##
     Min. 1st Ou. Median
                            Mean 3rd Ou.
                                             Max.
   0.9907 1.5061 1.6644 1.6576 1.8300 2.0264
  # Shortcut to get to the same place -
  summary(log(whr$happiness_score))
     Min. 1st Qu. Median
                            Mean 3rd Qu.
                                             Max.
   0.9907 1.5061 1.6644 1.6576 1.8300 2.0264
```



- This is a simple example of **metaprogramming** (that's the real name of this technique) and may seem trivial, but it's not
- For starters, you can't do it in Stata!



```
Copyright 1985-2017 StataCorp LLC
  Statistics/Data Analysis
                                      StataCorp
                                      4905 Lakeway Drive
     MP - Parallel Edition
                                      College Station, Texas 77845 USA
                                      800-STATA-PC
                                                          http://www.stata.com
                                      979-696-4600
                                                          stata@stata.com
                                      979-696-4601 (fax)
681-user 4-core Stata network perpetual license:
       Serial number: 501506002486
         Licensed to: WBG User
                       World Bank Group
Notes:
     1. Unicode is supported; see help unicode advice.
      2. More than 2 billion observations are allowed; see help obs advice.
      3. Maximum number of variables is set to 120000; see help set maxvar.
      4. New update available; type -update all-
running C:\Program Files (x86)\Statal5\sysprofile.do ...
. sysuse auto
(1978 Automobile Data)

    summarize log(make)

variable log not found
r(111);
```



- This is a **very powerful technique**, as you will soon see
- It's **also a common source of error**, as you can only use one function inside the other if the output of the inner function is the same as the input of the outer function
- It can also get quite tricky to follow what a line of code with multiple functions inceptions is doing
- Which is why we sometimes use pipes: %>%



```
# 1: Doing it the long way -----
# Create a vector with the log of the happiness score
log score <- log(whr$happy score)</pre>
# Get descriptive statistics for the log vector
mean(log score)
# 2: Shortcut to get to the same place ------
mean(log(whr$happy_score))
# 3: Now with pipes -----
whr$happy_score %>%
 log() %>%
 mean()
```

In a few words, x % f() is the same as f(x)



Now that you know piping exists in R, you should know that it can **drastically improve code readability**. And from now on you can also laugh if you see this in some tidyverse nerd laptop sticker or t-shirt:

%>%
magrittr

Ceci n'est pas un pipe.



#### Loops

- One thing that usually gives people away as Stata users writing R code are loops
- In Stata, we use for loops quite a lot
- The equivalent to that in R would be to write a for loop like this

```
# A for loop in R
for (number in 1:5) {
    print(number)
}
```



#### Loops

## [1] 3

## [1] 4 ## [1] 5

```
# A for loop in R
for (number in 1:5) {
    print(number)
}
## [1] 1
## [1] 2
```



#### **Apply**

- R, however, has a whole function family that allows users to loop through an object **in a more efficient way**, without using explicit loops
- They're called apply and there are many of them, with different use cases
- If you look for the apply help file, you can see all of them
- For the purpose of this training, we will only use two of them, sapply and apply



#### **Apply**

sapply(X, FUN, ...): applies a function to all elements of a vector or list and returns the result in a vector. Its arguments are:

- X: a data frame, matrix or vector the function will be applied to
- **FUN:** the function you want to apply
- ...: possible function arguments

## [1] 1 2



#### Loops vs Apply

```
# A for loop in R
for (number in c(1.2,2.5)) {
    print(round(number))
}

## [1] 1
## [1] 2

# A much more elegant loop in R
sapply(c(1.2,2.5), round)
```



#### Loops vs Apply

- When looping, you repeat the same operation over a set of items
- apply(), instead, takes all your elements at once and applies an operation to them simultaneously
- The difference is like this: imagine you ask a yes/no question to a group of people.
  - You can collect the answers by asking each one of them individually -- this is looping
  - Otherwise, you can ask them to raise their hands and collect the answers at once -- this is apply()



#### **Apply**

A more general version of sapply() is the apply() function.

apply(X, MARGIN, FUN, ...): applies a function to all columns or rows of matrix. Its arguments are

- X: a data frame (or matrix) the function will be applied to
- MARGIN: 1 to apply the function to all rows or 2 to apply the function to all columns
- **FUN:** the function you want to apply
- ...: possible function arguments



### **Apply**

```
## [,1] [,2] [,3]
## [1,] 1 6 2
## [2,] 24 9 74
## [3,] 9 4 2
```



### **Apply**

```
# Row means
apply(matrix, 1, mean)

## [1] 3.00000 35.66667 5.00000

# Column means
apply(matrix, 2, mean)

## [1] 11.333333 6.333333 26.000000
```



- As we have said several times, **R** is super flexible
- One example of that is that it's super easy and quick to create custom functions
- Here's how:



```
square <- function(x) {
    y <- x ^ 2
    return(y)
}
square(2)</pre>
```

```
## [1] 4
```



#### Exercise 9: Create a function

- 1. Create a function that calculates the z-score of a vector.
  - Recall the outline of functions in R:

```
function_name <- function(input) {
  output <- operation(input)
  return(output)
}</pre>
```

- Hints:
  - The command to obtain the mean of a vector is mean(x)
  - The command to get the SD of a vector is sd(x)
  - R is vectorized: you can operate vectors and number directly and the result will be a vector



```
zscore <- function(x) {
  mean <- mean(x, na.rm = T)
  sd <- sd(x, na.rm = T)
  z <- (x - mean)/sd
  return(z)
}
summary(zscore(whr$happiness_score))</pre>
```

```
## Min. 1st Qu. Median Mean 3rd Qu. Max.
## -2.3551 -0.7579 -0.0776 0.0000 0.7590 1.9492
```

## Indentation

#### Indentation

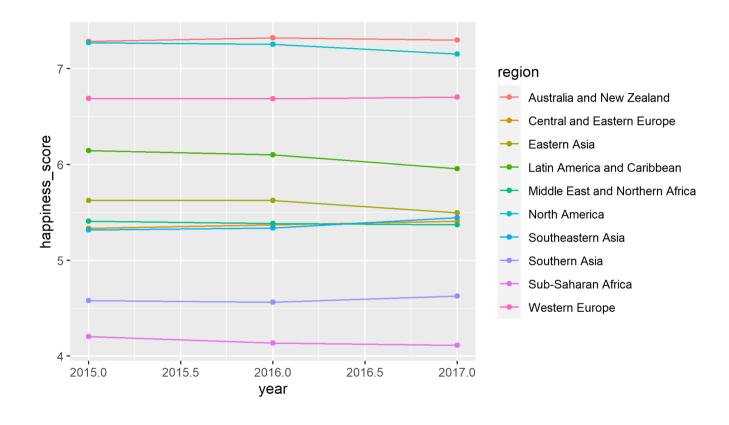


```
# Here's some code
annualHappy_reg <- aggregate(happy_score ~ year + region, data = whr, FUN = mean)
plot <- ggplot(annualHappy_reg,aes(y = happy_score,x = year, color = region, group = region))
+ geom_line() + geom_point()
print(plot)</pre>
```

```
# Here's the same code
annualHappy_reg <-
  aggregate(happiness_score ~ year + region,
            data = whr,
            FUN = mean)
plot <-
  ggplot(annualHappy_reg,
         aes(y = happiness_score,
             x = year,
             color = region,
             group = region)) +
  geom_line() +
  geom_point()
```

### Indentation







### Why indent?

- Even though R understands what unindented code says, it can be quite difficult for a human being to read it
- On the other hand, white space does not have a special meaning for R, so it will understand code that is more readable for a human being



### Why indent?

- Indentation in R looks different than in Stata:
  - To indent a whole line, you can select that line and press Tab
  - To unindent a whole line, you can select that line and press **Shift + Tab**
  - However, this will not always work for different parts of a code in the same line
- In R, we typically don't introduce white space manually
- It's rather introduced by RStudio for us



#### Exercise 8: Indentation in R

To see an example of how indenting works in RStudio, go back to our first example with sapply:

```
# A much more elegant loop in R sapply(c(1.2,2.5), round)
```

- 1. Add a line between the two arguments of the function (the vector of numbers and the round function)
- 2. Now add a line between the numbers in the vector.



Note that RStudio formats the different arguments of the function differently:

Thank you!

# Appendix

## Appendix - Initial settings



### .Rhistory and .RData

- .Rhistory automatically stores the commands entered in the console
- RData stores the objects in your environment only if you save your workspace, and loads them again in the next RStudio session
- Both files are relative to the working directory where your RStudio session started

## Appendix - Assignment 1



### Assignment 1

Create a function that

- 1. Takes as argument a vector of packages names
- 2. Loops through the packages listed in the input vector
- 3. Install the packages
- 4. Loads the packages

## Appendix - If statements



#### If statements

- Installing packages can be time-consuming, especially as the number of packages you're using grows, and each package only needs to be installed once
- We often use locals in Stata to create section switches to install packages
- In R, the equivalent to that would be to create a new object as a section switch

#### Exercise 9: Creating an if statement

Create a dummy scalar object called PACKAGES.

• TIP: Section switches can also be Boolean objects.

## Appendix - If statements



- Now we need to create an if statement using this switch
- If statements in R look like this:

## Appendix - If statements



#### If statements

Possible variations would include

```
# Turn switch on
PACKAGES <- TRUE

# Using a Boolean object

if (PACKAGES == TRUE) {
   install.packages(packages, dep = TRUE)
}

# Which is the same as

if (PACKAGES) {
   install.packages(packages, dep = TRUE)
}</pre>
```

## Appendix - Assignment 2



#### Create a function that

- 1. Takes as argument a vector of packages names
- 2. Loops through the packages listed in the input vector
- 3. Tests if a package is already installed
- 4. Only installs packages that are not yet installed
- 5. Loads the packages
- TIP: to test if a package is already installed, use the following code:

```
# Test if object x is contained in
# the vector of installed packages
x %in% installed.packages()
```

## Appendix - File paths best practices



### File paths best practices

- We at DIME Analytics recommend always using **explicit** and **dynamic** file paths
- **Explicit** means you're explicitly stating where the file will be saved -- instead of setting the working directory, for example
- **Dynamic** means that you don't need to adjust every file path in the script when you change from one machine to another -- they're updated based on a single line of code to be changed

## Appendix - File paths best practices



### File paths

• Explicit and dynamic file path:

## Appendix - Using packages



### Using packages

Once a package is loaded, you can use its features and functions. Here's a list of some useful and cool packages:

- Rcmdr easy to use GUI
- swirl an interactive learning environment for R and statistics.
- ggplot2 beautiful and versatile graphics (the syntax is a pain, though)
- stargazer awesome latex regression and summary statistics tables
- foreign reads .dta and other formats from inferior statistical software
- **zoo** time series and panel data manipulation useful functions
- data.table some functions to deal with huge data sets
- sp and rgeos spatial analysis
- multiwayvcov and sandwich clustered and robust standard errors
- RODBC, RMySQL, RPostgresSQL, RSQLite For relational databases and using SQL in R.

## Appendix - Resources



#### Resources

- A discussion of folder structure and data managament can be found here: https://dimewiki.worldbank.org/wiki/DataWork\_Folder
- For a broader discussion of data management, go to https://dimewiki.worldbank.org/wiki/Data\_Management

## Appendix - Git



### Git

Git is a version-control system for tracking changes in code and other text files. It is a great resource to include in your work flow.

We didn't cover it here because of time constraints, but below are some useful links, and DIME Analytics provides trainings on Git and GitHub, so keep an eye out for them.

- DIME Analytics git page: https://worldbank.github.io/dimeanalytics/git/
- A Quick Introduction to Version Control with Git and GitHub: https://journals.plos.org/ploscompbiol/article? id=10.1371/journal.pcbi.1004668

## Appendix - R projects



### R projects

If you have used R before, you may have heard of RStudio Projects. It's RStudio's suggested tool for workflow management. DIME Analytics has found that it is not the best fit for our needs, because

- 1. In DIME, we mainly use Stata, and we prefer to keep a similar structure in R (Stata 15 also has a projects feature, but it is not yet widely adopted)
- 2. We need to keep our code and data in separate folders, as we store code in GitHub and data in DropBox

However, if you want to learn more about it, we recommend starting here: https://r4ds.had.co.nz/workflow-projects.html